

# Mining Massive Datasets

B. Ravindran

**Reconfigurable and Intelligent Systems Engineering (RISE) Group**  
**Department of Computer Science and Engineering**

**Robert Bosch Centre for Data Science and Artificial Intelligence (RBC-DSAI)**  
**Indian Institute of Technology Madras**



# Ensembles

# Problem Setting

$$X^{(1)} = \langle 0.15, 0.25 \rangle, Y^{(1)} = -1$$

$$X^{(2)} = \langle 0.4, 0.45 \rangle, Y^{(2)} = +1$$

$$\vdots$$

The input can be thought of as  $X = (X_1, X_2, \dots, X_p)$

The  $X_i$  are the features that describe the input.

The output is either a categorical variable, typically denoted by  $G$   
or a continuous variable denoted by  $Y$

The  $i$ -th instance of the input is denoted as  $x_i$  and output is denoted as  $y_i$

We loosely state the learning problem as given a value of input  $X$   
make a good prediction  $\hat{Y}$  of  $Y$  or  $\hat{G}$  of  $G$

# Problem Setting

$\mathcal{X} \subseteq \mathfrak{R}^p$  is the input space

$X = (X_1, X_2, \dots, X_p)$  is a random variable describing the input

$\mathcal{Y} \subseteq \mathfrak{R}$  or  $\Gamma$  is the output space

$Y$  is a random variable describing the output


$p(X, Y)$  is the data distribution

$$p(X, Y) = p(Y|X)p(X)$$

$p(Y|x)$  is the predicted output probabilities given an input  $x$

# Ensemble Methods: Increasing the Accuracy

- Ensemble methods
  - Use a combination of models to increase accuracy
  - Combine a series of  $k$  learned models,  $M_1, M_2, \dots, M_k$ , with the aim of creating an improved model  $M^*$
- Popular ensemble methods
  - Bagging: averaging the prediction over a collection of classifiers
  - Boosting: weighted vote with a collection of classifiers
  - Ensemble: combining a set of heterogeneous classifiers



# Bagging: Bootstrap Aggregation

- Training
  - Given a set  $D$  of  $d$  tuples, at each iteration  $i$ , a training set  $D_i$  of  $d$  tuples is sampled with replacement from  $D$  (i.e., bootstrap)
  - A classifier model  $M_i$  is learned for each training set  $D_i$
- Classification: classify an unknown sample  $X$ 
  - Each classifier  $M_i$  returns its class prediction
  - The bagged classifier  $M^*$  counts the votes and assigns the class with the most votes to  $X$
- Accuracy
  - Often significantly better than a single classifier derived from  $D$
  - For noisy data: not considerably worse, more robust
  - Proved improved accuracy in prediction
- More *Stable*



# Bagging: Example





# Random Forests

## Bagging Decision Trees



1. If the number of cases in the training set is  $N$ , sample  $N$  cases at random - but *with replacement*, from the original data. This sample will be the training set for growing the tree.
2. If there are  $M$  input variables, a number  $m \ll M$  is specified such that at each node,  $m$  variables are selected at random out of the  $M$  and the best split on these  $m$  is used to split the node. The value of  $m$  is held constant during the forest growing.
3. Each tree is grown to the largest extent possible. There is no pruning.





# Why Random Forests Work?

- The *correlation* between any two trees in the forest needs to be low. **Increasing the correlation increases the forest error rate.**
- A tree with a low error rate is a *strong* classifier. **Increasing the strength of the individual trees decreases the forest error rate.**
- Reducing  $m$  reduces both the correlation and the strength. Increasing it increases both.
- Somewhere in between is an "optimal" range of  $m$  - usually quite wide.



# Features of Random Forest



- Excellent accuracy among current algorithms.
- Runs efficiently on large data bases.
- Can handle thousands of input variables without variable deletion.
- Gives estimates of what variables are important in the classification.
- Generates an internal unbiased estimate of the generalization error as the forest building progresses.

# Error Estimates

- Out-of-Bag Error Estimates
  - We know that about 27% of the data is not sampled in bootstrap
  - This is true for each tree
  - Hence each data point is not used in about a third of the trees.
- Measure average error on each data point from the trees not using them
  - All these trees vote on the class of this data point
  - If the predicted majority class doesn't match the true class then error
- Shown to be unbiased sample of the error

# Committee Methods

- Takes a simple unweighted average of the predictions from each model
- Assigns equal probability to each model.
- Applicable in cases where the different models arise from the same parametric model, with different parameter values.

# Stacking

- Learning methods are “stacked” on top of one another .
- Combines multiple models’ output with estimated optimal weights.
- Leads to better prediction.
- Train a “predictor of predictors”
  - Treat individual predictors as features
  - Similar to multi-layer perceptron idea
  - Special case: binary,  $f_e$  linear  $\Rightarrow$  weighted vote

$$\hat{y}_1 = f_1(x_1, x_2, \dots)$$

$$\hat{y}_2 = f_2(x_1, x_2, \dots) \quad \Rightarrow \quad \hat{y}_e = f_e(\hat{y}_1, \hat{y}_2, \dots)$$

...

# Boosting

- Focus new learners on examples that others get wrong
- Train learners sequentially
- Errors of early predictions indicate the “hard” examples
- Focus later predictions on getting these examples right
- Combine the whole set in the end
- Convert many “weak” learners into a complex predictor

# Boosting

- How boosting works?
  - Weights are assigned to each training tuple
  - A series of  $k$  classifiers is iteratively learned
  - After a classifier  $M_i$  is learned, the weights are updated to allow the subsequent classifier,  $M_{i+1}$ , to pay more attention to the training tuples that were misclassified by  $M_i$
  - The final  $M^*$  combines the votes of each individual classifier, where the weight of each classifier's vote is a function of its accuracy
- Comparing with bagging: boosting tends to achieve greater accuracy, but it also risks overfitting the model to misclassified data
  - Can be shown to maximize margin of classifier

# AdaBoost (Freund and Schapire, 1997)

- The current linear combination of classifiers is

$$C_{(m-1)}(x_i) = \alpha_1 k_1(x_i) + \alpha_2 k_2(x_i) + \cdots + \alpha_{m-1} k_{m-1}(x_i)$$

- Extend it to,

$$C_m(x_i) = C_{(m-1)}(x_i) + \alpha_m k_m(x_i)$$

- Total cost, or total error, of the extended classifier as the exponential loss

$$E = \sum_{i=1}^N e^{-y_i(C_{(m-1)}(x_i) + \alpha_m k_m(x_i))}$$

Source: AdaBoost and the Super Bowl of Classifiers  
A Tutorial Introduction to Adaptive Boosting - Raúl Rojas



- Since our intention is to draft  $k_m$  we rewrite the above expression as,

$$E = \sum_{i=1}^N w_i^{(m)} e^{-y_i \alpha_m k_m(x_i)}$$

where

$$w_i^{(m)} = e^{-y_i C_{(m-1)}(x_i)}$$

- Split the sum into two sums

$$E = \sum_{y_i = k_m(x_i)} w_i^{(m)} e^{-\alpha_m} + \sum_{y_i \neq k_m(x_i)} w_i^{(m)} e^{\alpha_m}$$

- Simplify the notation to

$$E = W_c e^{-\alpha_m} + W_e e^{\alpha_m}$$

Source: AdaBoost and the Super Bowl of Classifiers  
A Tutorial Introduction to Adaptive Boosting - Raúl Rojas

# AdaBoost - Weighting

- To determine weight of  $m^{\text{th}}$  classifier,

$$\frac{dE}{d\alpha_m} = -W_c e^{-\alpha_m} + W_e e^{\alpha_m}$$

- Equating it to zero,

$$\alpha_m = \frac{1}{2} \ln \left( \frac{W_c}{W_e} \right)$$

- Rewriting, with  $W$  as the total sum of weights,

$$\alpha_m = \frac{1}{2} \ln \left( \frac{W - W_e}{W_e} \right) = \frac{1}{2} \ln \left( \frac{1 - e_m}{e_m} \right) \quad \text{Where } e_m = W_e / W$$

Source: AdaBoost and the Super Bowl of Classifiers  
A Tutorial Introduction to Adaptive Boosting - Raúl Rojas

# AdaBoost - Algorithm

## AdaBoost

For  $m = 1$  to  $M$

1. Select and extract from the pool of classifiers the classifier  $k_m$  which minimizes

$$W_e = \sum_{y_i \neq k_m(x_i)} w_i^{(m)}$$

2. Set the weight  $\alpha_m$  of the classifier to

$$\alpha_m = \frac{1}{2} \ln \left( \frac{1 - e_m}{e_m} \right)$$

where  $e_m = W_e / W$

3. Update the weights of the data points for the next iteration. If  $k_m(x_i)$  is a miss, set

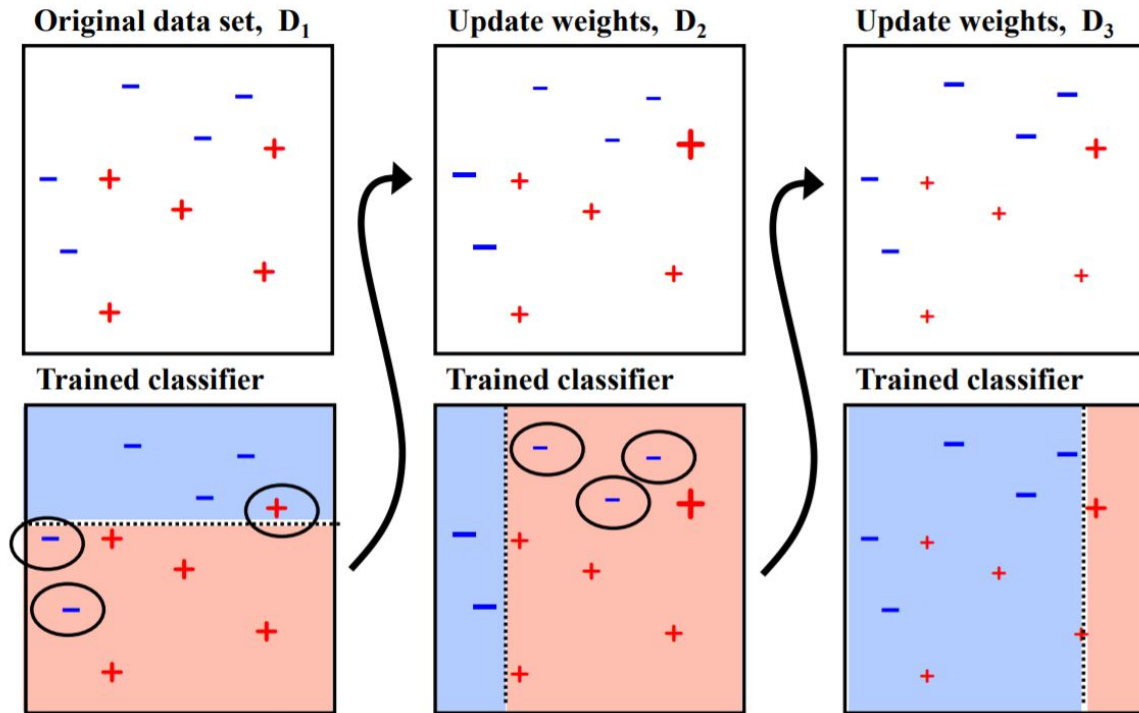
$$w_i^{(m+1)} = w_i^{(m)} e^{\alpha_m} = w_i^{(m)} \sqrt{\frac{1 - e_m}{e_m}}$$

otherwise

$$w_i^{(m+1)} = w_i^{(m)} e^{-\alpha_m} = w_i^{(m)} \sqrt{\frac{e_m}{1 - e_m}}$$

Credits : ESL


# AdaBoost - Example




Source : Slides from Dr.Kalev Kask's (UCI) lecture on Ensembles of Learners

# AdaBoost - Example

Weight each classifier and combine them:




Combined classifier




Source : Slides from Dr.Kalev Kask's  
(UCI) lecture on Ensembles of Learners

# Adaboost



*Simulated data test error rate for boosting with stumps, as a function of the number of iterations. Also shown are the test error rate for a single stump, and a 244-node classification tree.*



*Simulated data, boosting with stumps*

Credits : ESL




# Gradient Boosting

- Learn sequence of predictors
- Subsequent models predict the error residual of the previous predictions
- Sum of predictions is increasingly accurate
- Predictive function is increasingly complex

# Gradient Boosting - Example


Learn a simple predictor...

$$f_1(x^{(i)}) \approx y^{(i)}$$



Then try to correct its errors

$$\epsilon^{(i)} = y^{(i)} - f_1(x^{(i)})$$




Source : Slides from Dr.Kalev Kask's  
(UCI) lecture on Ensembles of Learners



# Gradient Boosting - Example


Combining gives a better predictor...

$$\Rightarrow f_1(x^{(i)}) + f_2(x^{(i)}) \approx y^{(i)}$$



Can try to correct its errors also, & repeat

$$\epsilon_2^{(i)} = y^{(i)} - f_1(x^{(i)}) - f_2(x^{(i)}) \quad \dots$$




Source : Slides from Dr.Kalev Kask's  
(UCI) lecture on Ensembles of Learners


# Gradient Boosting - Example

$$y^{(i)} \approx \sum_z f_z(x^{(i)})$$

Data & prediction function



Error residual



Source : Slides from Dr.Kalev Kask's  
(UCI) lecture on Ensembles of Learners

# Gradient Boosted Trees

---

**Algorithm 10.3** *Gradient Tree Boosting Algorithm.*

---

1. Initialize  $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^N L(y_i, \gamma)$ .
2. For  $m = 1$  to  $M$ :

(a) For  $i = 1, 2, \dots, N$  compute

$$r_{im} = - \left[ \frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}.$$

(b) Fit a regression tree to the targets  $r_{im}$  giving terminal regions  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$ .

(c) For  $j = 1, 2, \dots, J_m$  compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma).$$

(d) Update  $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$ .

3. Output  $\hat{f}(x) = f_M(x)$ .
- 

Credits : ESL

# Summary

- Ensembles yield powerful classifiers
- Random forest is one of the best performing classifiers
- Bagging results in more stable classifiers
  - Can be parallelized
- Boosting gives better performance
  - Can overfit
- Gradient Boosted Decision Trees are very powerful